

# Multilayer Perceptron

ITM528 Deep learning

Taemoon Jeong

# Multilayer Perceptron

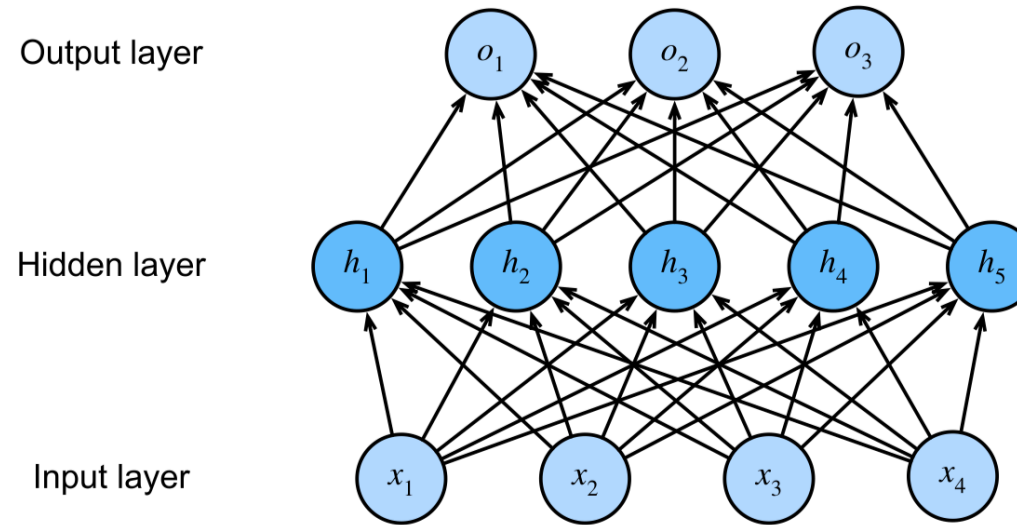
# Multilayer Perceptrons

- We move on to truly (and the simplest) deep neural networks called **multilayer perceptrons** (MLPs).
  - They consist of multiple layers of neurons fully connected to those in the layer below and those above.
- Although automatic differentiation significantly simplifies the implementation, we will see how the gradients are calculated in deep networks.
- MLPs are high-capacity models and they are often prone to overfitting.
  - Thus, we will revisit regularization and generalization for deep networks.

# Multilayer Perceptrons

- Our previous linear classification model simply maps inputs directly to outputs via a single affine transformation followed by a softmax operation.
- However, **linearity** (in affine transformation) is a strong assumption.
  - For example, linearity implies the weaker assumption of monotonicity.
  - However, most real-world problems are **nonlinear** (e.g., health assessments as a function of body temperature or dog classifier).
- One way to overcome this is to find a suitable **representation** of inputs, on top of which a single linear model would suffice.

# Hidden Layers



- We can overcome the limitations of linear models by incorporating **one or more hidden layers**.
  - This architecture is commonly called a multilayer perceptron, aka **MLP**.

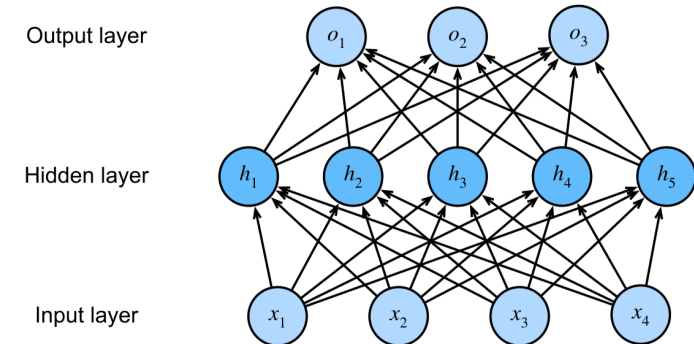
# From Linear to Nonlinear

- Suppose we denote the matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$  a minibatch of  $n$  examples with  $d$  inputs (features).
- For a one-hidden-layer MLP whose hidden layer has  $h$  hidden units, we denote  $\mathbf{H} \in \mathbb{R}^{n \times h}$  the outputs of the hidden layer (hidden representations).
- Then, the output of the MLP is calculated by:

$$\mathbf{H} = \mathbf{XW}^{(1)} + \mathbf{b}^{(1)}$$

$$\mathbf{O} = \mathbf{HW}^{(2)} + \mathbf{b}^{(2)}$$

- However, simply adding the hidden layer does not change anything!
  - Why?

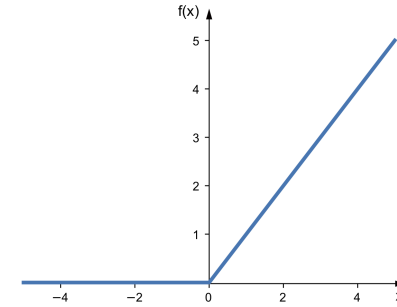


$$\mathbf{O} = (\mathbf{XW}^{(1)} + \mathbf{b}^{(1)})\mathbf{W}^{(2)} + \mathbf{b}^{(2)} = \mathbf{XW} + \mathbf{b}$$

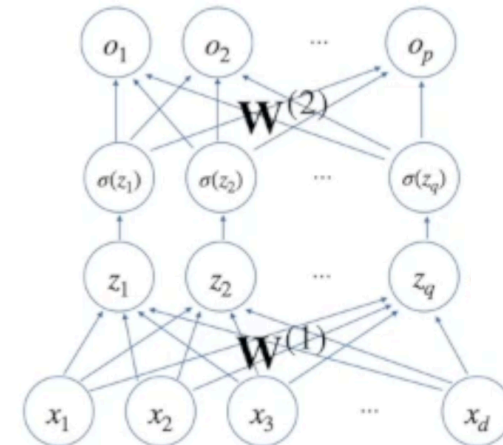
# From Linear to Nonlinear

- Hence, we need a **nonlinear activation** function  $\sigma$  to be applied to each hidden unit following the affine transformation.
  - One popular choice is the ReLU (Rectified Linear Unit),  $\sigma(x) = \max(0, x)$ .
- Combining all together,

$$\mathbf{H} = \sigma(\mathbf{XW}^{(1)} + \mathbf{b}^{(1)})$$
$$\mathbf{O} = \mathbf{HW}^{(2)} + \mathbf{b}^{(2)}$$



- Of course, we can build more general MLPs by continuously stacking such hidden layers.



# Universal Approximators

- It is worth asking just **how powerful** a deep network could be.
- A universal approximation property states that
- George Cybenko proved that an **MLP** with one hidden layer and sigmoid activations has the **universal approximation** property [1].
  - Polynomial functions also satisfy this property [2]. In fact, **polynomial functions** are the first function classes to have this property:

## Fundamental Theorem of Approximation Theory

Let  $f \in C[a, b]$ ,  $-\infty < a < b < \infty$ . Given  $\epsilon > 0$ , there exists an algebraic polynomial  $p$  for which

$$|f(x) - p(x)| < \epsilon$$

for all  $x \in [a, b]$ .

- In fact, if one can show that other families of functions (e.g., deep networks) behave like polynomials, then such families also have universal approximation properties.
- **Kernel methods** also satisfy this property [3].

[1] "Approximation by superpositions of a sigmoidal function," 1989

[2] "A generalized Weierstrass approximation theorem," 1948

[3] "Universal Approximation Using Radial-Basis-Function Networks," 1991

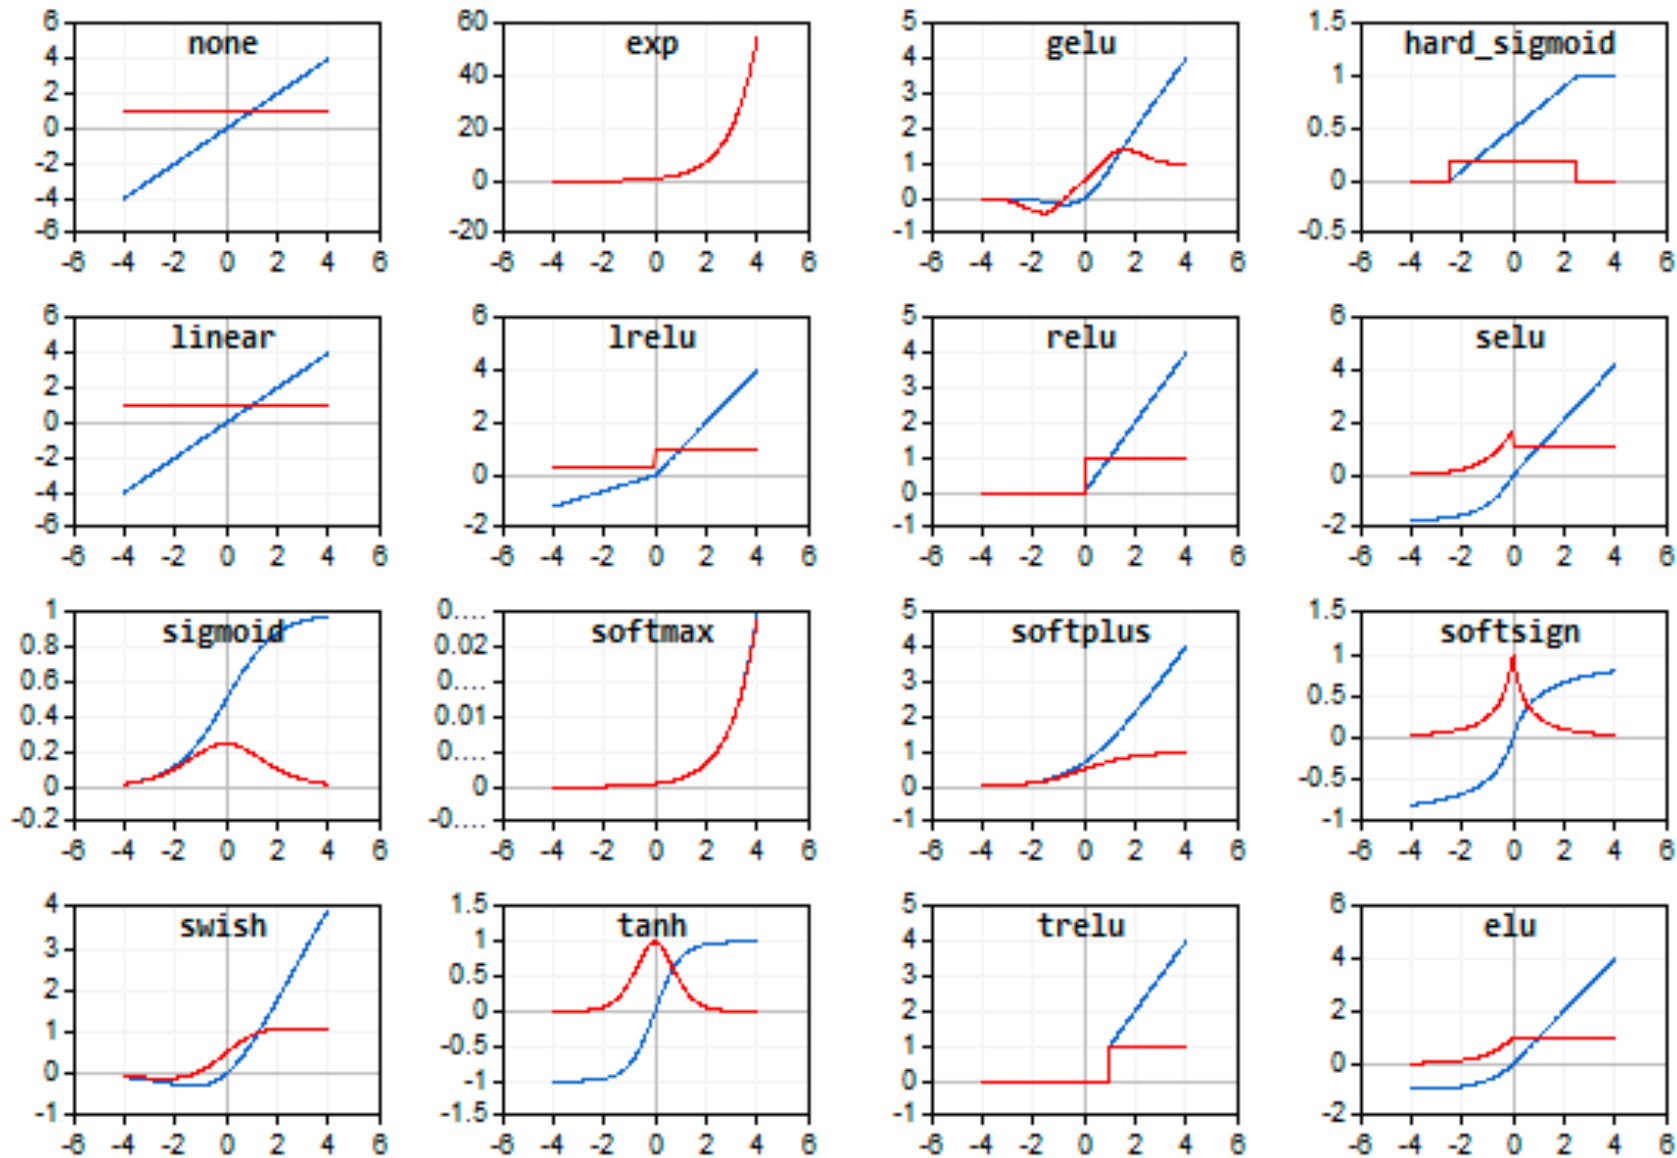


# Universal Approximators



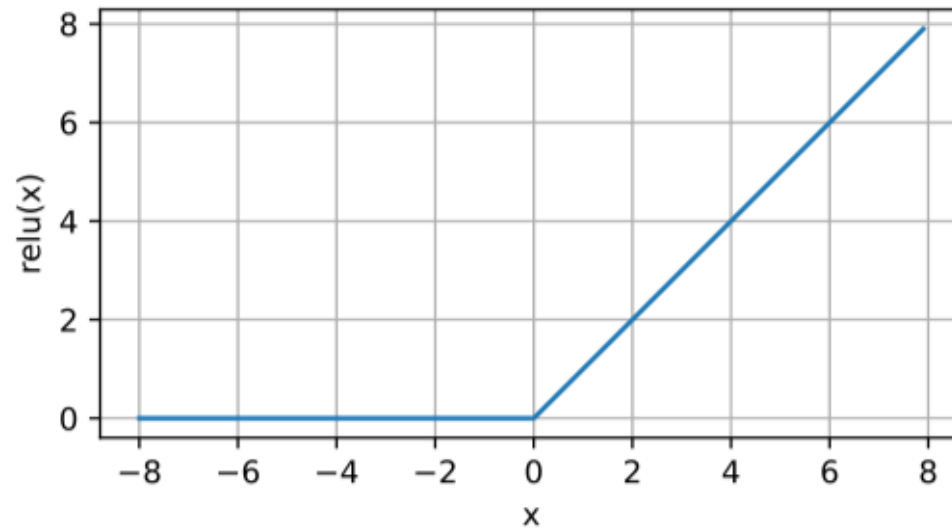
- Note that these results only suggest that a single-hidden-layer network with enough nodes can model any continuous function.
  - Hence, just because a single-hidden-layer network can learn any function, it does not mean that you should try to solve all of your problems with this.

# Activation Functions



# Activation Functions

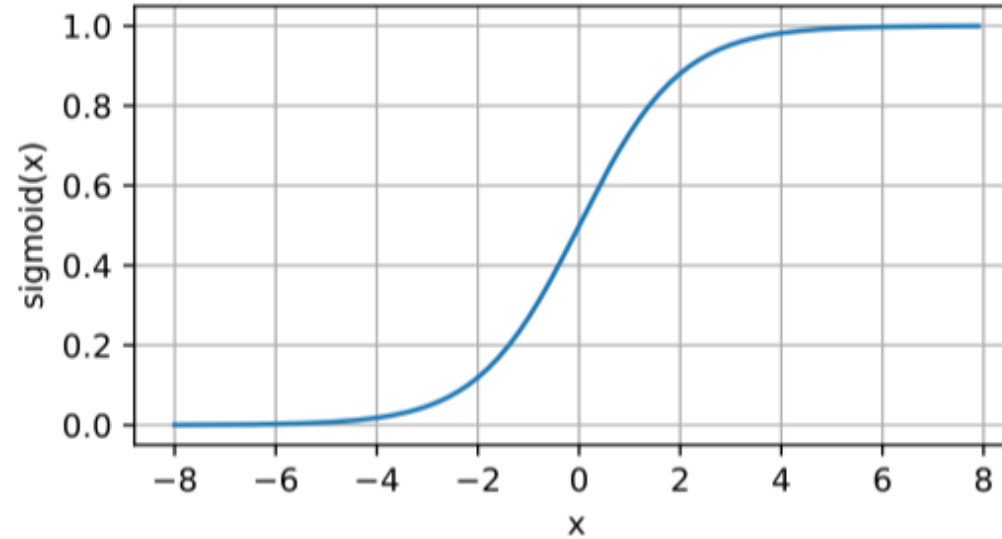
- Activation functions are differentiable operators to transform input signals to outputs to add **non-linearity**.
- ReLU (rectified linear unit) Function:  $\text{ReLU}(x) = \max(x, 0)$



- A ReLU activation function is one of the most popular activation functions in deep learning and has its strength in handling the **vanishing gradient issue**.
- However, a dead relu problem might occur.

# Activation Functions

• Sigmoid Function:  $\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$

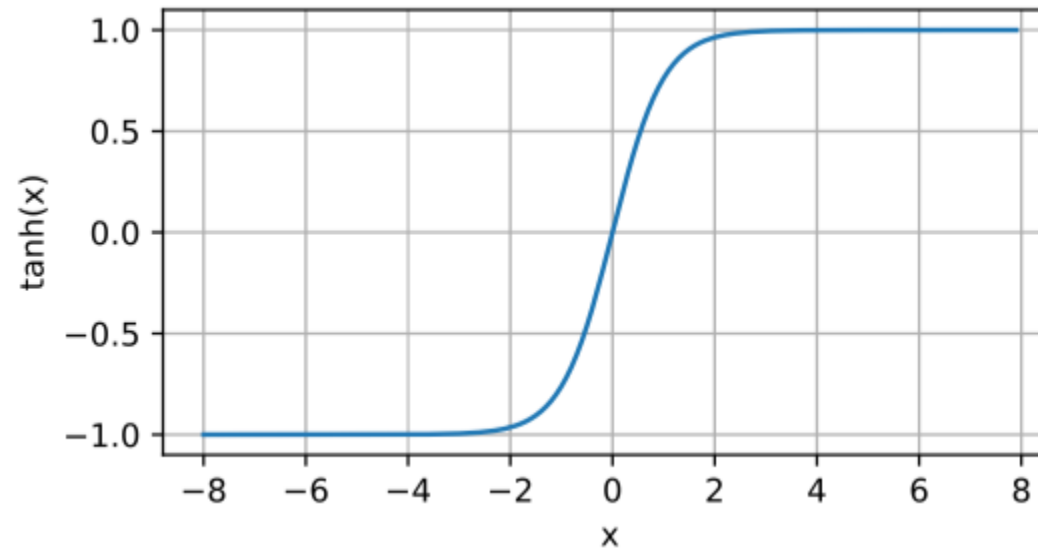


- The sigmoid function squashes the output to lie on the interval (0,1).
- The gradient vanishing problem might happen, hindering its performance on deep networks.
- Also, the outputs are not zero-centered.
- One useful property of the sigmoid is that

$$\frac{d}{dx}\text{sigmoid}(x) = \text{sigmoid}(x)(1 - \text{sigmoid}(x))$$

# Activation Functions

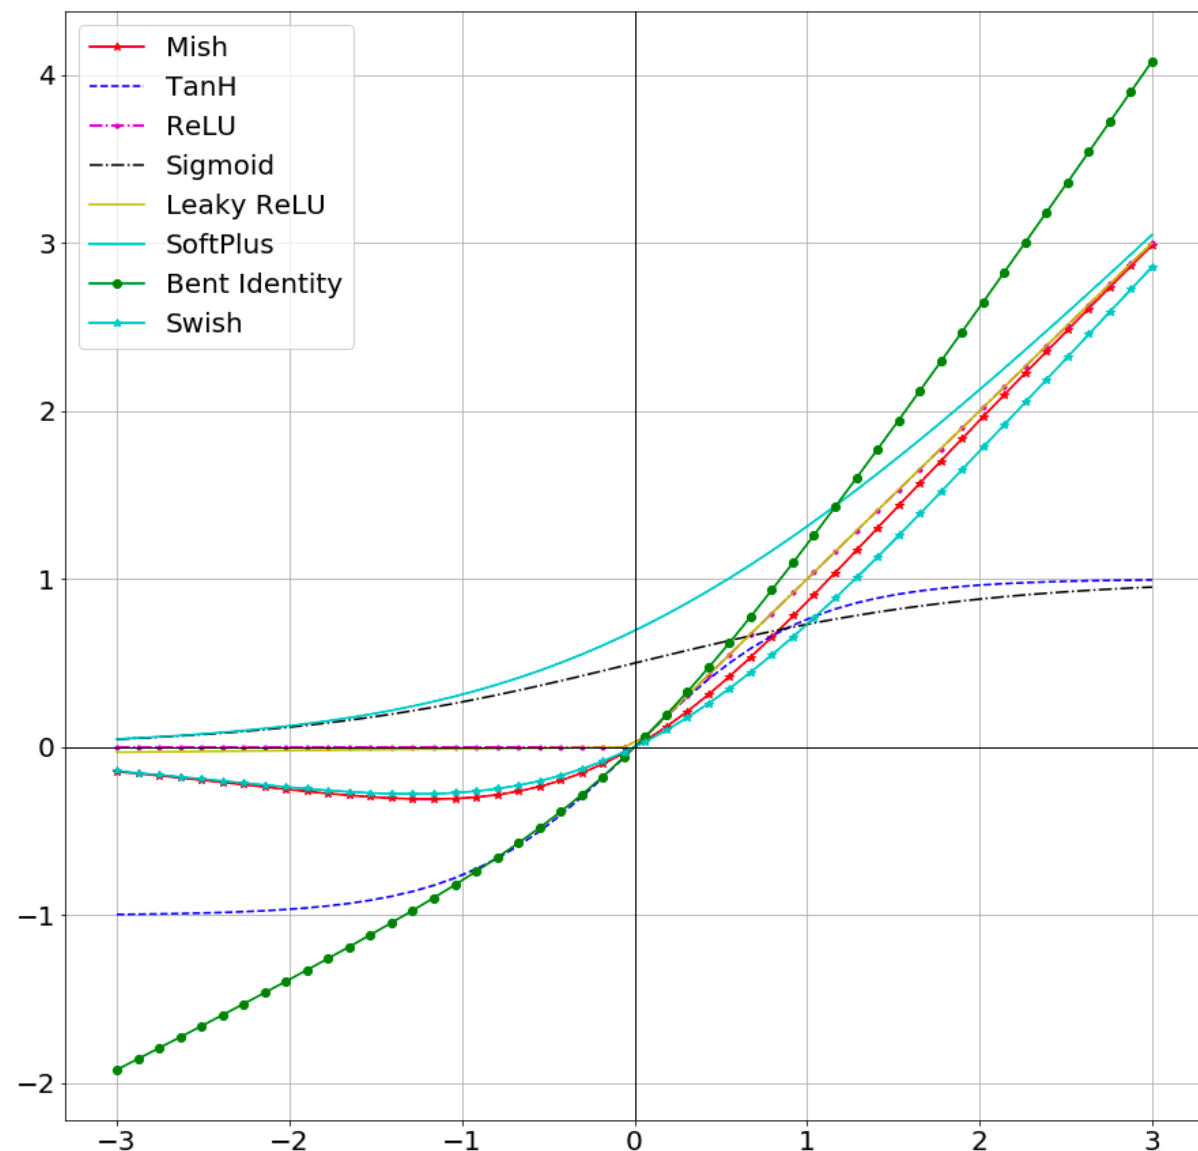
• Tanh Function:  $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$



- The tanh function squashes the output to lie on the interval  $(-1,1)$ , hence symmetric.
- The vanishing gradient problem still exists.
- One useful property of the tanh is that

$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$

# Activation Functions



- The **relu** is notorious for the dead relu problem.
- To handle this, the **elu** function was proposed. However, it introduces a longer computation time due to the exponential operation included.

$$\text{elu}(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha(e^x - 1), & x < 0 \end{cases}$$

- The **leaky relu** function also avoids the dead relu problem and is fast. However, we have to tune the parameter  $\alpha$ .

$$\text{lrelu}(x) = \begin{cases} x, & \text{if } x > 0 \\ \alpha x, & x < 0 \end{cases}$$

- The **gelu** function works well in NLP, specifically Transformer models, as it is fast.

$$\text{gelu}(x) = 0.5x(1 + \tanh(\sqrt{2/\pi}(x + 0.044715x^3)))$$

- The **swish** function is continuous and differentiable at all points. And it works well on standard image datasets (CIFAR or ImageNet) compared to others (relu, lrelu, elu, gelu).

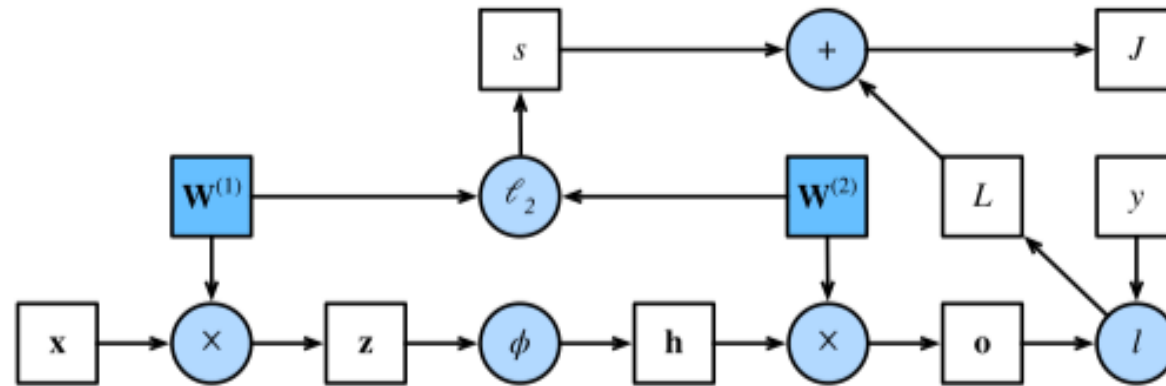
$$\text{swish}(x) = x(1 + e^{-x})^{-1}$$

- The **mish** function is  $C^\infty$ -continuous and approximates identity near the origin. In some experiments, the **mish** works better than **swish** activations.

$$\text{mish}(x) = x \tanh(\text{softplus}(x))$$

# Backpropagation

# Forward Propagation



- [1] For the sake of simplicity, an input example is  $\mathbf{x} \in \mathbb{R}^d$  and no bias term exists.
- [2] Then, the intermediate variable is  $\mathbf{z} = \mathbf{W}^{(1)}\mathbf{x} \in \mathbb{R}^h$  where  $\mathbf{W}^{(1)} \in \mathbb{R}^{h \times d}$  is the weight parameter of the hidden layer.
- [3] Our hidden activation vector is  $\mathbf{h} = \phi(\mathbf{z}) \in \mathbb{R}^h$ .
- [4] The hidden layer output becomes  $\mathbf{o} = \mathbf{W}^{(2)}\mathbf{h} \in \mathbb{R}^q$  where  $\mathbf{W}^{(2)} \in \mathbb{R}^{q \times h}$  is the weight parameter of the hidden layer.
- [5] The loss term for a single data becomes  $L = l(\mathbf{o}, y)$ .
- [6] Also, the regularization term becomes  $s = \frac{\lambda}{2}(\|\mathbf{W}^{(1)}\|_F^2 + \|\mathbf{W}^{(2)}\|_F^2)$ .
- [7] Finally, the model's regularized loss on a given data example is:  $J = L + s$ .

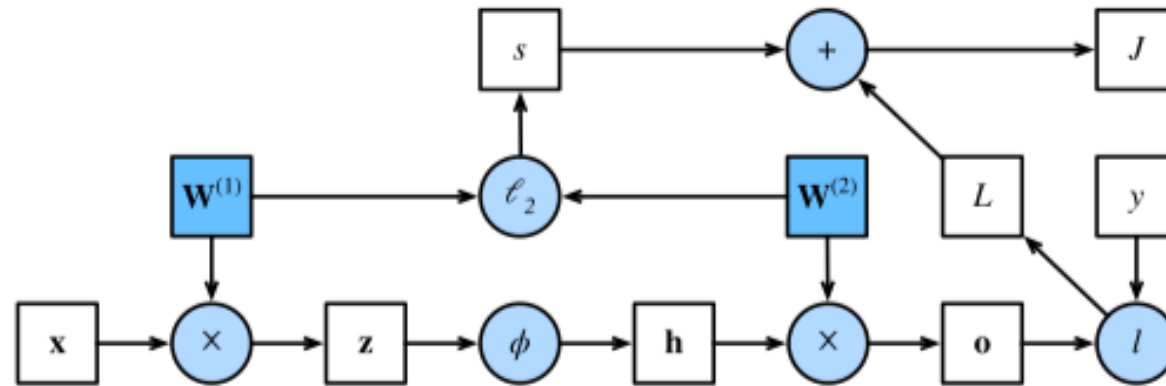


# Backpropagation

- **Backpropagation** refers to the method of calculating the gradient of neural network parameters.
- In short, this method traverses the network in reverse order, from the output to the input layer, according to the **chain rule** from calculus.
  - Assume that we have  $Y = f(X)$  and  $Z = g(Y)$ .
  - By using the chain rule:

$$\frac{\partial Z}{\partial X} = \prod \left( \frac{\partial Z}{\partial Y}, \frac{\partial Y}{\partial X} \right)$$

# Backpropagation



$$\mathbf{o} = \mathbf{W}^{(2)}\mathbf{h} \in \mathbb{R}^q$$

$$s = \frac{\lambda}{2}(\|\mathbf{W}^{(1)}\|_F^2 + \|\mathbf{W}^{(2)}\|_F^2)$$

- [1] The first step is to calculate the gradients of the objective  $J = L + s$  w.r.t. the loss term  $L$  and the regularization term  $s$ .

$$\frac{\partial J}{\partial L} = 1 \text{ and } \frac{\partial J}{\partial s} = 1$$

- [2] Next, we compute the gradient of  $J$  w.r.t. the output layer  $\mathbf{o}$ :

$$\frac{\partial J}{\partial \mathbf{o}} = \Pi\left(\frac{\partial J}{\partial L}, \frac{\partial L}{\partial \mathbf{o}}\right) = \frac{\partial L}{\partial \mathbf{o}} \in \mathbb{R}^q$$

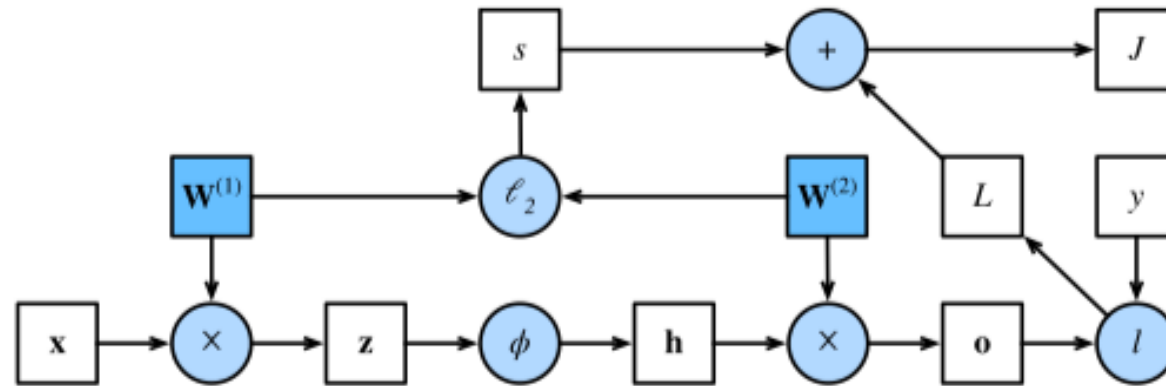
- [3] Next, we calculate the gradients of  $s$  w.r.t. parameters  $\mathbf{W}^{(1)}$  and  $\mathbf{W}^{(2)}$ :

$$\frac{\partial s}{\partial \mathbf{W}^{(1)}} = \lambda \mathbf{W}^{(1)} \text{ and } \frac{\partial s}{\partial \mathbf{W}^{(2)}} = \lambda \mathbf{W}^{(2)}$$

- [4] Now, we are able to calculate the gradient  $\partial J / \partial \mathbf{W}^{(2)} \in \mathbb{R}^{q \times h}$ :

$$\frac{\partial J}{\partial \mathbf{W}^{(2)}} = \Pi\left(\frac{\partial J}{\partial \mathbf{o}}, \frac{\partial \mathbf{o}}{\partial \mathbf{W}^{(2)}}\right) + \Pi\left(\frac{\partial J}{\partial s}, \frac{\partial s}{\partial \mathbf{W}^{(2)}}\right) = \frac{\partial J}{\partial \mathbf{o}} \mathbf{h}^T + \lambda \mathbf{W}^{(2)}$$

# Backpropagation



$$\mathbf{o} = \mathbf{W}^{(2)}\mathbf{h} \in \mathbb{R}^q$$

$$s = \frac{\lambda}{2}(\|\mathbf{W}^{(1)}\|_F^2 + \|\mathbf{W}^{(2)}\|_F^2)$$

- [1] To obtain the gradient w.r.t.  $\mathbf{W}^{(1)}$ , we need to continue backpropagation along the output layer to the hidden layer

$$\frac{\partial J}{\partial \mathbf{h}} = \prod \left( \frac{\partial J}{\partial \mathbf{o}}, \frac{\partial \mathbf{o}}{\partial \mathbf{h}} \right) = \mathbf{W}^{(2)T} \frac{\partial J}{\partial \mathbf{o}}$$

- [2] Since the activation function  $\phi$  applies elementwise:

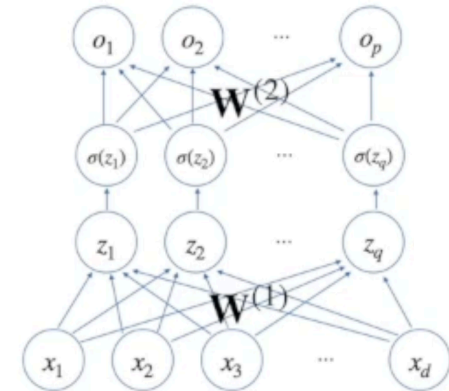
$$\frac{\partial J}{\partial \mathbf{z}} = \prod \left( \frac{\partial J}{\partial \mathbf{h}}, \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right) = \frac{\partial J}{\partial \mathbf{h}} \odot \phi'(\mathbf{z})$$

- [3] Finally, we can obtain the gradient  $\partial J / \partial \mathbf{W}^{(1)} \in \mathbb{R}^{h \times d}$ :

$$\frac{\partial J}{\partial \mathbf{W}^{(1)}} = \prod \left( \frac{\partial J}{\partial \mathbf{z}}, \frac{\partial \mathbf{z}}{\partial \mathbf{W}^{(1)}} \right) + \prod \left( \frac{\partial J}{\partial s}, \frac{\partial s}{\partial \mathbf{W}^{(1)}} \right) = \frac{\partial J}{\partial \mathbf{z}} \mathbf{x}^T + \lambda \mathbf{W}^{(1)}$$

- [4] We can further express this with:

$$\frac{\partial J}{\partial \mathbf{W}^{(1)}} = \left( \left( \mathbf{W}^{(2)T} \frac{\partial L}{\partial \mathbf{o}} \right) \odot \phi'(\mathbf{W}^{(1)}\mathbf{x}) \right) \mathbf{x}^T + \lambda \mathbf{W}^{(1)}$$



# Matrix Calculus

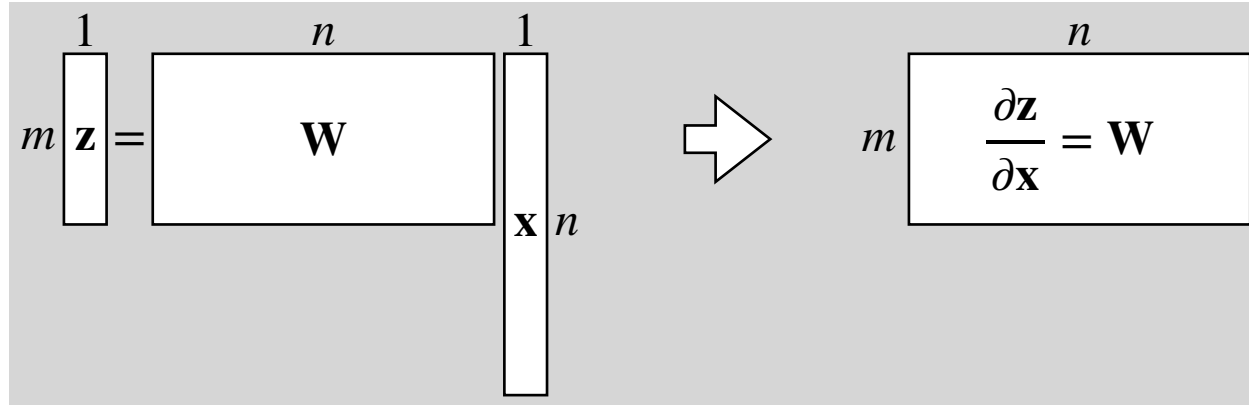
$$x \in \mathbb{R}^n \mapsto f \in \mathbb{R}^m \quad \begin{array}{c} 1 \\ \vdots \\ n \end{array} \longrightarrow \begin{array}{c} 1 \\ \vdots \\ m \end{array} \Rightarrow \frac{\partial f}{\partial x} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix} \quad \begin{array}{c} n \\ \frac{\partial f}{\partial x} \in \mathbb{R}^{m \times n} \end{array}$$

$$x \in \mathbb{R}^n \mapsto f \in \mathbb{R}^m \mapsto g \in \mathbb{R}^k \quad \begin{array}{c} 1 \\ \vdots \\ n \end{array} \longrightarrow \begin{array}{c} 1 \\ \vdots \\ m \end{array} \longrightarrow \begin{array}{c} 1 \\ \vdots \\ k \end{array} \Rightarrow \frac{\partial g}{\partial x} = \frac{\partial g}{\partial f} \frac{\partial f}{\partial x}$$

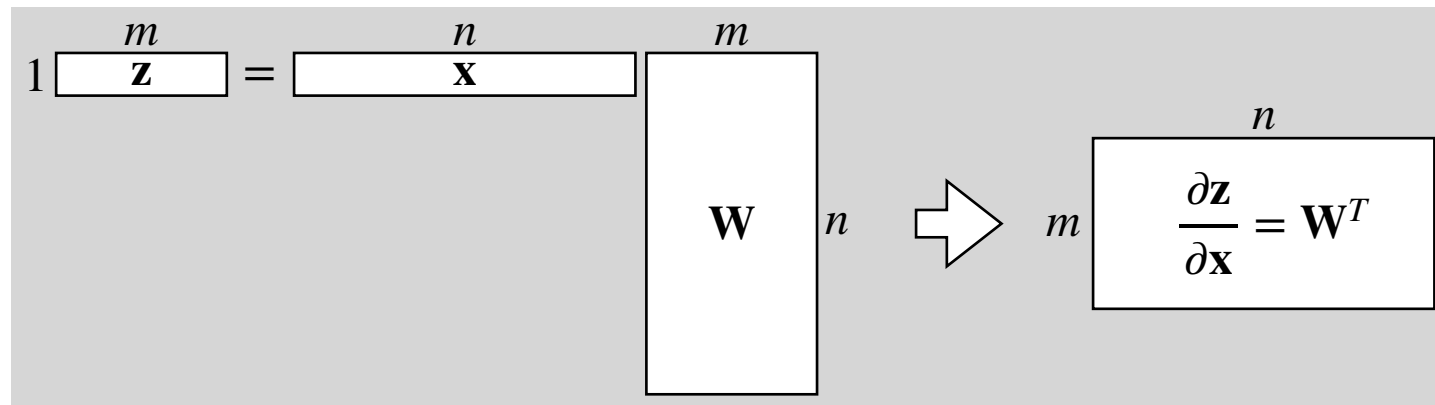
$$\begin{array}{c} \frac{\partial g}{\partial x} \in \mathbb{R}^{k \times n} \\ k \quad \begin{array}{c} \phantom{\frac{\partial g}{\partial x}} \\ n \end{array} \end{array} = \begin{array}{c} \frac{\partial g}{\partial f} \in \mathbb{R}^{k \times m} \\ k \quad \begin{array}{c} \phantom{\frac{\partial g}{\partial f}} \\ m \end{array} \end{array} \begin{array}{c} \frac{\partial f}{\partial x} \in \mathbb{R}^{m \times n} \\ \begin{array}{c} \phantom{\frac{\partial f}{\partial x}} \\ n \end{array} \end{array}$$

# Matrix Calculus

- Matrix times column vector:  $\mathbf{z} = \mathbf{W}\mathbf{x}$  where  $\mathbf{z} \in \mathbb{R}^{m \times 1}$ ,  $\mathbf{x} \in \mathbb{R}^{n \times 1}$ , and  $\mathbf{W} \in \mathbb{R}^{m \times n}$ . Then,  $\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \mathbf{W} \in \mathbb{R}^{m \times n}$ .

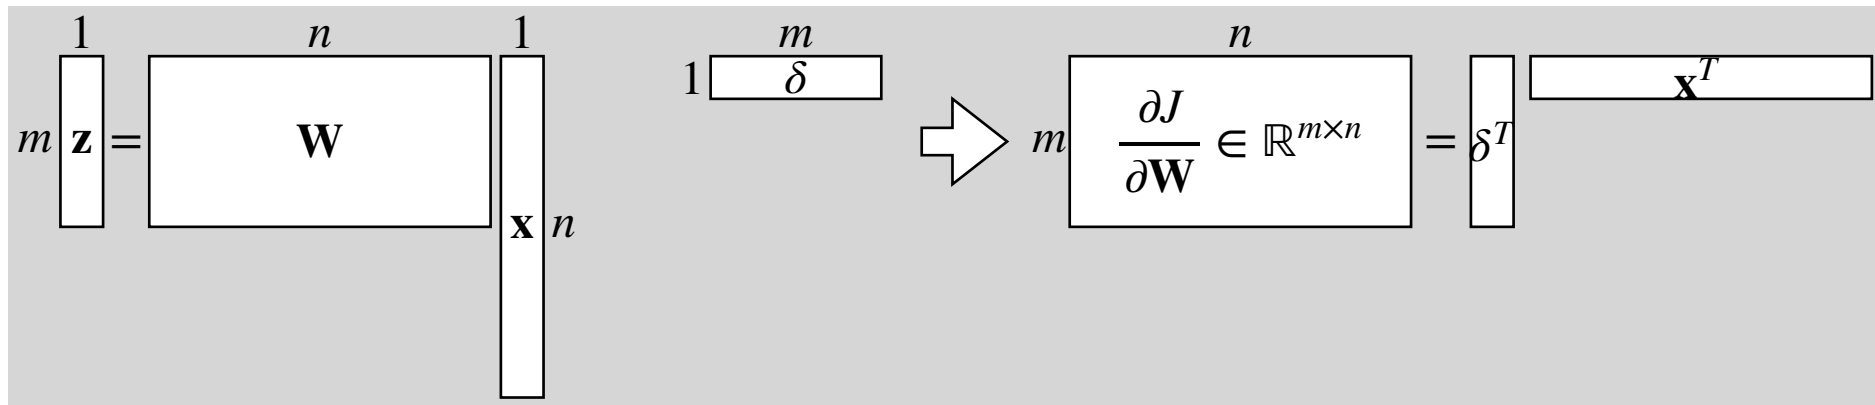


- Row vector times matrix:  $\mathbf{z} = \mathbf{x}\mathbf{W}$  where  $\mathbf{z} \in \mathbb{R}^{1 \times m}$ ,  $\mathbf{x} \in \mathbb{R}^{1 \times n}$ , and  $\mathbf{W} \in \mathbb{R}^{n \times m}$ . Then,  $\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \mathbf{W}^T \in \mathbb{R}^{m \times n}$ .



# Matrix Calculus

- A vector with itself:  $\mathbf{z} = \mathbf{x}$ . Then,  $\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \mathbf{I}$ .
- An elementwise function applied to a vector:  $\mathbf{z} = f(\mathbf{x})$  where  $z_i = f(x_i)$ . Then,  $\frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \text{diag}(f'(\mathbf{x}))$ .
  - We can also write  $\odot f'(\mathbf{x})$  when applying the chain rule!
- Matrix times column vector:  $\mathbf{z} = \mathbf{W}\mathbf{x}$  and  $\delta = \frac{\partial J}{\partial \mathbf{z}} \in \mathbb{R}^{1 \times m}$ . Then,  $\frac{\partial J}{\partial \mathbf{W}} = \frac{\partial J}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \delta^T \mathbf{x}^T \in \mathbb{R}^{m \times n}$



- Row vector times Matrix:  $\mathbf{x} \in \mathbb{R}^{1 \times n}$ ,  $\mathbf{W} \in \mathbb{R}^{n \times m}$ ,  $\mathbf{z} = \mathbf{x}\mathbf{W} \in \mathbb{R}^{1 \times m}$ , and  $\delta = \frac{\partial J}{\partial \mathbf{z}} \in \mathbb{R}^{1 \times m}$ .

• Then,  $\frac{\partial J}{\partial \mathbf{W}} = \frac{\partial J}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \delta \frac{\partial \mathbf{z}}{\partial \mathbf{W}} = \mathbf{x}^T \delta \in \mathbb{R}^{n \times m}$

# Matrix Calculus

- (Recall) Cross-entropy loss w.r.t. logits:
  - Suppose  $J = \text{CE}(\mathbf{y}, \hat{\mathbf{y}}) \in \mathbb{R}$  and  $\hat{\mathbf{y}} = \text{softmax}(\mathbf{z}) \in \mathbb{R}^{1 \times k}$ .
    - Then,  $\frac{\partial J}{\partial \mathbf{z}} = \hat{\mathbf{y}} - \mathbf{y} \in \mathbb{R}^{1 \times k}$  where  $k$  is the number of classes.

# Vanishing and Exploding Gradients

- Consider a deep network with  $L$  layers, input  $\mathbf{x}$ , and output  $\mathbf{o}$ .
- With each layer  $l$  defined by a transformation  $f_l$  parametrized by weights  $\mathbf{W}^{(l)}$ , whose hidden layer output is  $\mathbf{h}^{(l)}$  (let  $\mathbf{h}^{(0)} = \mathbf{x}$ ), our network can be expressed as:

$$\mathbf{h}^{(l)} = f_l(\mathbf{h}^{(l-1)}) \text{ and thus } \mathbf{o} = f_L \circ \cdots \circ f_1(\mathbf{x})$$

- If all the hidden layer output and the input are vectors, we can write the gradient of  $\mathbf{o}$  with respect to  $\mathbf{W}^{(l)}$ :

$$\partial_{\mathbf{W}^{(l)}} \mathbf{o} = \underbrace{\partial_{\mathbf{h}^{(L-1)}} \mathbf{h}^{(L)}}_{=\mathbf{M}^{(L)}} \cdots \underbrace{\partial_{\mathbf{h}^{(l)}} \mathbf{h}^{(l+1)}}_{=\mathbf{M}^{(l+1)}} \underbrace{\partial_{\mathbf{W}^{(l)}} \mathbf{h}^{(l)}}_{=\mathbf{v}^{(l)}}$$

- In other words, this gradient is the **product of  $L - 1$  matrices** and the gradient vector.



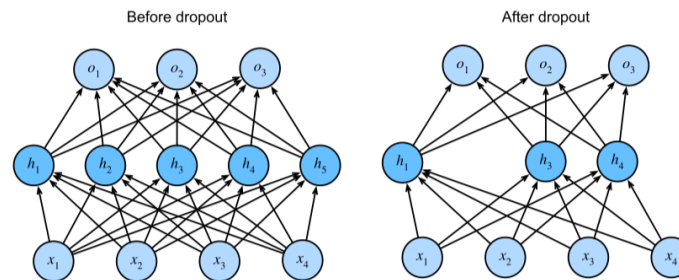
# Other Issues

# Nonparametrics?

- Are MLPs parametric models?
  - The models do have millions of parameters.
- While neural networks clearly have parameters, in some ways, it can be more fruitful to think of them as behaving like nonparametric models.
- So what precisely makes a model nonparametric?
  - While the name covers a diverse set of approaches, one common theme is that nonparametric methods tend to have a level of complexity that grows as the amount of available data grows.
  - Nonparametric methods include 1) k-nearest neighbor algorithm and 2) kernel-based methods.
- In a sense, because neural networks are over-parametrized, they tend to interpolate the training data (fitting it perfectly) having the same property as nonparametric models.

# Regularization in Neural Networks

- Early stopping
  - While deep neural networks are capable of fitting arbitrary labels, early stopping becomes an efficient method for regularizing networks.
  - In other words, whenever a model has fitted the cleanly labeled data but not randomly labeled examples, it has, in fact, been generalized.
- Weight decay
  - Depending on which weight norm is penalized, it is known either as ridge regularization (for  $l_2$ -norm) or lasso regularization (for  $l_1$ -norm).
  - While it still remains a popular tool, researchers have noted that typical strengths of weight decay are insufficient for generalization.
- Dropout
  - Randomly replace some portion of nodes into 0.

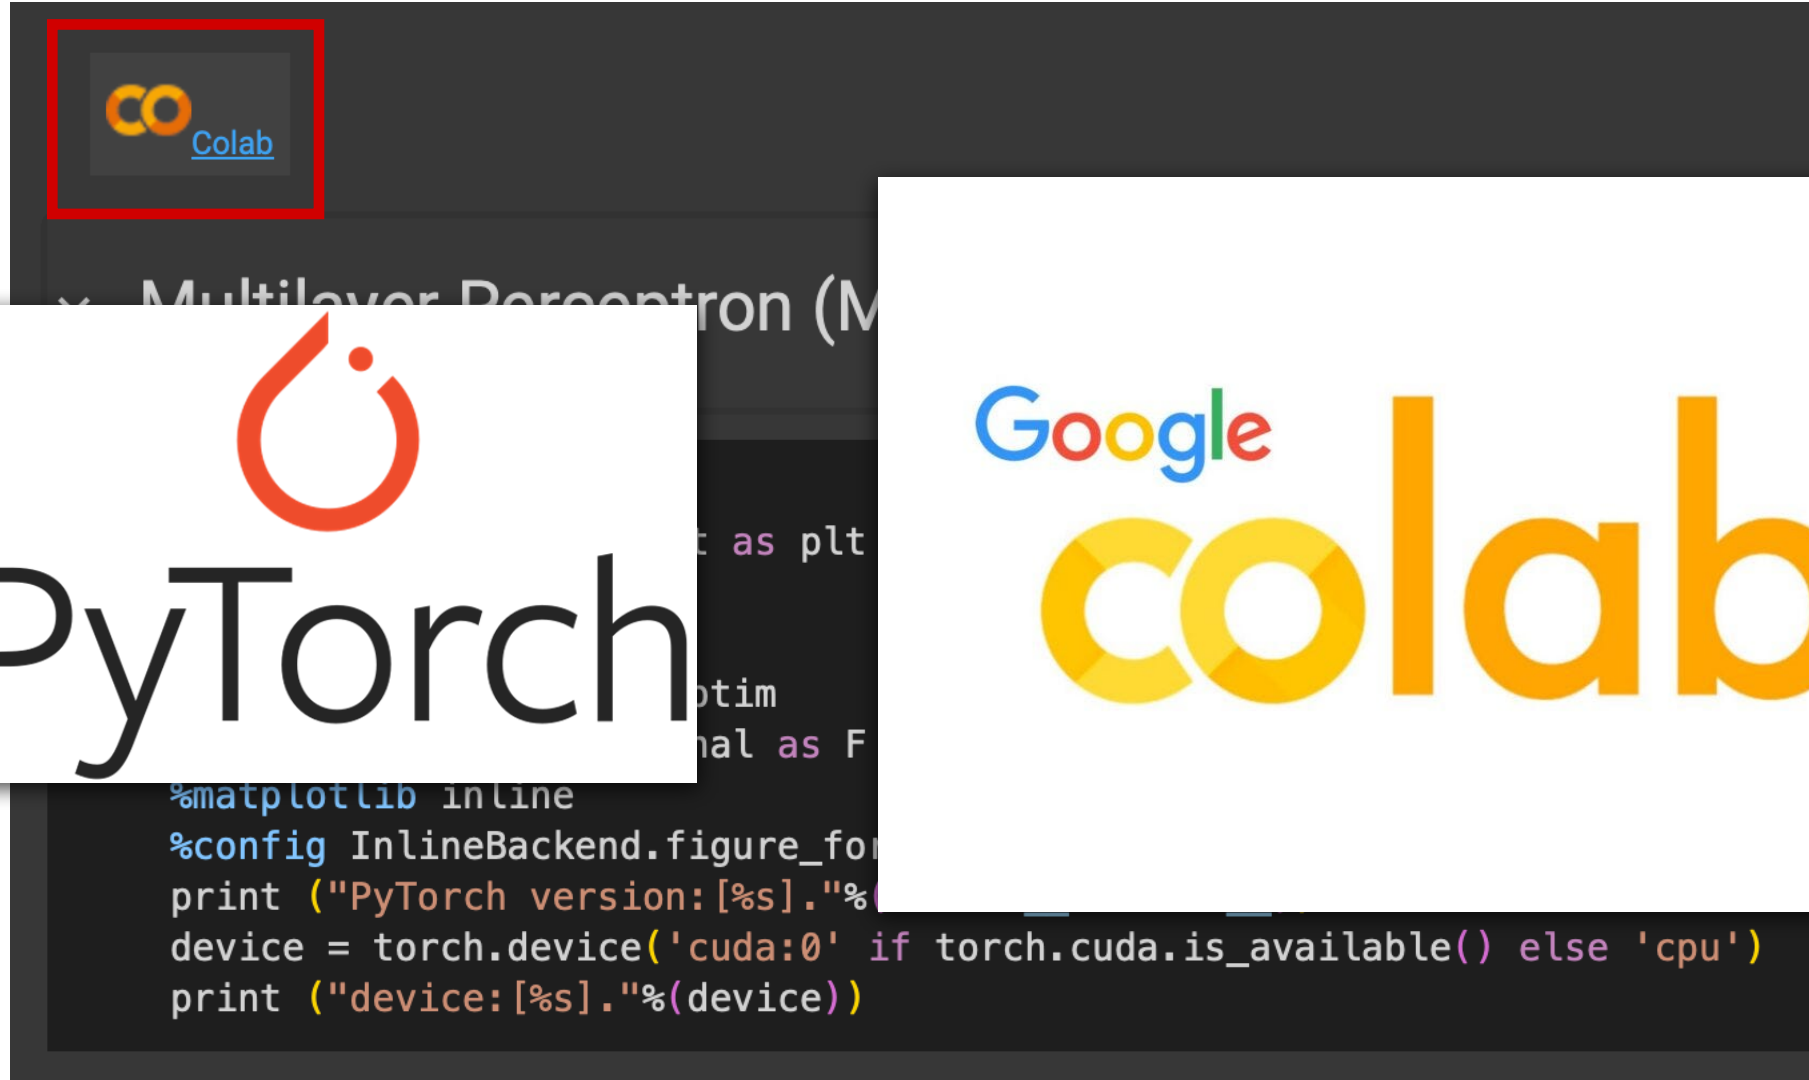


# Code Implementation (MLP)

Colab Link:

<https://colab.research.google.com/drive/13JTcJ1mBCk5ZZH68yRb3LqmnxTLrz4Ua?authuser=2#scrollTo=uaokkwJwsN5I>

# 1. Library Import



# 1. Library Import



## ✓ Multilayer Perceptron (MLP)

```
[ ] import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
%matplotlib inline
%config InlineBackend.figure_format='retina'
print ("PyTorch version:[%s]."%(torch.__version__))
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print ("device:[%s]."%(device))
```

Importing Libraries

## 2. Load the dataset (MNIST)

### Dataset (MNIST)



```

▶ from torchvision import datasets, transforms
mnist_train = datasets.MNIST(root='./data')
mnist_test = datasets.MNIST(root='./data')
print ("mnist_train:\n",mnist_train,"\n")
print ("mnist_test:\n",mnist_test,"\n")
print ("Done.")

```

ue)  
ue)

## 2. Load the dataset (MNIST)

### Dataset (MNIST)



```
from torchvision import datasets, transforms
mnist_train = datasets.MNIST(root='./data/', train=True, transform=transforms.ToTensor(), download=True)
mnist_test = datasets.MNIST(root='./data/', train=False, transform=transforms.ToTensor(), download=True)
print ("mnist_train:\n",mnist_train,"\n")
print ("mnist_test:\n",mnist_test,"\n")
print ("Done.")
```

Loading the MNIST Dataset



# 3. Batch size and Data Iterator

## ▼ Data Iterator

### 1. Setting Batch size

```
[ ] BATCH_SIZE = 256
train_iter = torch.utils.data.DataLoader(mnist_train, batch_size=BATCH_SIZE, shuffle=True, num_workers=1)
test_iter = torch.utils.data.DataLoader(mnist_test, batch_size=BATCH_SIZE, shuffle=True, num_workers=1)
print ("Done.")
```

### 2. Creating Data Loaders

# 4. Define the MLP model



## Define the MLP model

```
class MultiLayerPerceptronClass(nn.Module):
    """
    Multilayer Perceptron (MLP) Class
    """
    def __init__(self, name='mlp', xdim=784, hdim=256, ydim=10):
        super(MultiLayerPerceptronClass, self).__init__()
        self.name = name
        self.xdim = xdim
        self.hdim = hdim
        self.ydim = ydim
        self.lin_1 = nn.Linear(
            # FILL IN HERE
            self.xdim, self.hdim
        )
        self.lin_2 = nn.Linear(
            # FILL IN HERE
            self.hdim, self.ydim
        )
        self.init_param() # initialize parameters
```

```
def init_param(self):
    nn.init.kaiming_normal_(self.lin_1.weight)
    nn.init.zeros_(self.lin_1.bias)
    nn.init.kaiming_normal_(self.lin_2.weight)
    nn.init.zeros_(self.lin_2.bias)
```

```
def forward(self, x):
    net = x
    net = self.lin_1(net)
    net = F.relu(net)
    net = self.lin_2(net)
    return net
```

```
M = MultiLayerPerceptronClass(name='mlp', xdim=784, hdim=256, ydim=10).to(device)
loss = nn.CrossEntropyLoss()
optm = optim.Adam(M.parameters(), lr=1e-3)
print ("Done.")
```

# 4. Define the MLP model

## ▼ Define the MLP model



```
class MultiLayerPerceptronClass(nn.Module):  
    """  
    Multilayer Perceptron (MLP) Class  
    """  
    def __init__(self, name='mlp', xdim=784, hdim=256, ydim=10):  
        super(MultiLayerPerceptronClass, self).__init__()  
        self.name = name  
        self.xdim = xdim  
        self.hdim = hdim  
        self.ydim = ydim  
        self.lin_1 = nn.Linear(  
            # FILL IN HERE  
            self.xdim, self.hdim  
        )  
        self.lin_2 = nn.Linear(  
            # FILL IN HERE  
            self.hdim, self.ydim  
        )  
        self.init_param() # initialize parameters
```

Class Definition

# 4. Define the MLP model

## ▽ Define the MLP model



```
class MultiLayerPerceptronClass(nn.Module):
```

```
    """
```

```
        Multilayer Perceptron (MLP) Class
```

```
    """
```

```
    def __init__(self, name='mlp', xdim=784, hdim=256, ydim=10):
```

```
        super(MultiLayerPerceptronClass, self).__init__()
```

```
        self.name = name
```

```
        self.xdim = xdim
```

```
        self.hdim = hdim
```

```
        self.ydim = ydim
```

```
        self.lin_1 = nn.Linear(
```

```
            # FILL IN HERE
```

```
            self.xdim, self.hdim
```

```
        )
```

```
        self.lin_2 = nn.Linear(
```

```
            # FILL IN HERE
```

```
            self.hdim, self.ydim
```

```
        )
```

```
        self.init_param() # initialize parameters
```

initial setting

# 4. Define the MLP model

## ▽ Define the MLP model

```
class MultiLayerPerceptronClass(nn.Module):  
    """  
    Multilayer Perceptron (MLP) Class  
    """  
    def __init__(self, name='mlp', xdim=784, hdim=256, ydim=10):  
        super(MultiLayerPerceptronClass, self).__init__()  
        self.name = name  
        self.xdim = xdim  
        self.hdim = hdim  
        self.ydim = ydim  
        self.lin_1 = nn.Linear(  
            # FILL IN HERE  
            self.xdim, self.hdim  
        )  
        self.lin_2 = nn.Linear(  
            # FILL IN HERE  
            self.hdim, self.ydim  
        )  
        self.init_param() # initialize parameters
```

Linear Layers

# 4. Define the MLP model

## ✓ Define the MLP model

```
def init_param(self):  
    nn.init.kaiming_normal_(self.lin_1.weight)  
    nn.init.zeros_(self.lin_1.bias)  
    nn.init.kaiming_normal_(self.lin_2.weight)  
    nn.init.zeros_(self.lin_2.bias)
```

```
def forward(self,x):  
    net = x  
    net = self.lin_1(net)  
    net = F.relu(net)  
    net = self.lin_2(net)  
    return net
```

Parameter Initialization

```
M = MultiLayerPerceptronClass(name='mlp',xdim=784,hdim=256,ydim=10).to(device)  
loss = nn.CrossEntropyLoss()  
optm = optim.Adam(M.parameters(),lr=1e-3)  
print ("Done.")
```

```
,  
    self.init_param() # initialize parameters
```

# 4. Define the MLP model

## ✓ Define the MLP model

```
def init_param(self):  
    nn.init.kaiming_normal_(self.lin_1.weight)  
    nn.init.zeros_(self.lin_1.bias)  
    nn.init.kaiming_normal_(self.lin_2.weight)  
    nn.init.zeros_(self.lin_2.bias)
```

```
def forward(self,x):  
    net = x  
    net = self.lin_1(net)  
    net = F.relu(net)  
    net = self.lin_2(net)  
    return net
```

Forward propagation

```
M = MultiLayerPerceptronClass(name='mlp',xdim=784,hdim=256,ydim=10).to(device)  
loss = nn.CrossEntropyLoss()  
optm = optim.Adam(M.parameters(),lr=1e-3)  
print ("Done.")
```

```
,  
    self.init_param() # initialize parameters
```

# 4. Define the MLP model

## ✓ Define the MLP model

```
def init_param(self):  
    nn.init.kaiming_normal_(self.lin_1.weight)  
    nn.init.zeros_(self.lin_1.bias)  
    nn.init.kaiming_normal_(self.lin_2.weight)  
    nn.init.zeros_(self.lin_2.bias)
```

```
def forward(self,x):  
    net = x  
    net = self.lin_1(net)  
    net = F.relu(net)  
    net = self.lin_2(net)  
    return net
```

MLP Model, Loss Function, and Optimizer Setup

```
M = MultiLayerPerceptronClass(name='mlp',xdim=784,hdim=256,ydim=10).to(device)  
loss = nn.CrossEntropyLoss()  
optm = optim.Adam(M.parameters(),lr=1e-3)  
print ("Done.")
```

```
,  
    self.init_param() # initialize parameters
```



# 5. Evaluation Function

## ✓ Evaluation Function

```
[ ] def func_eval(model,data_iter,device):  
    with torch.no_grad():  
        model.eval() # evaluate (affects Dropout and BN)  
        n_total,n_correct = 0,0  
        for batch_in,batch_out in data_iter:  
            y_trgt = batch_out.to(device)  
            model_pred = model(  
                # FILL IN HERE  
                batch_in.view(-1, 28*28).to(device)  
            )  
            _,y_pred = torch.max(model_pred.data,1)  
            n_correct += (  
                # FILL IN HERE  
                y_pred == y_trgt  
            ).sum().item()  
            n_total += batch_in.size(0)  
        val_accr = (n_correct/n_total)  
        model.train() # back to train mode  
    return val_accr  
print ("Done")
```

# 5. Evaluation Function

## ✓ Evaluation Function

Switching to Evaluation Mode

```
[ ] def func eval(model,data_iter,device):  
    with torch.no_grad():  
        model.eval() # evaluate (affects Dropout and BN)  
        n_total,n_correct = 0,0  
        for batch_in,batch_out in data_iter:  
            y_trgt = batch_out.to(device)  
            model_pred = model(  
                # FILL IN HERE  
                batch_in.view(-1, 28*28).to(device)  
            )  
            _,y_pred = torch.max(model_pred.data,1)  
            n_correct += (  
                # FILL IN HERE  
                y_pred == y_trgt  
            ).sum().item()  
            n_total += batch_in.size(0)  
        val_accr = (n_correct/n_total)  
        model.train() # back to train mode  
    return val_accr  
print ("Done")
```

# 5. Evaluation Function

## ✓ Evaluation Function

```
[ ] def func_eval(model,data_iter,device):  
    with torch.no_grad():  
        model.eval() # evaluate (affects Dropout and BN)  
        n_total,n_correct = 0,0  
        for batch_in,batch_out in data_iter:  
            y_trgt = batch_out.to(device)  
            model_pred = model(  
                # FILL IN HERE  
                batch_in.view(-1, 28*28).to(device)  
            )  
            y_pred = torch.max(model_pred.data,1)  
            # FILL IN HERE  
            y_pred == y_trgt  
            ).sum().item()  
            n_total += batch_in.size(0)  
            val_accr = (n_correct/n_total)  
            model.train() # back to train mode  
        return val_accr  
    print ("Done")
```

Looping Through Data and Making Predictions

# 5. Evaluation Function

## ✓ Evaluation Function

```
[ ] def func_eval(model,data_iter,device):  
    with torch.no_grad():  
        model.eval() # evaluate (affects Dropout and BN)  
        n_total,n_correct = 0,0  
        for batch_in,batch_out in data_iter:  
            y_trgt = batch_out.to(device)  
            model_pred = model(  
                # FILL IN HERE  
            )  
            _,y_pred = torch.max(model_pred.data,1)  
            n_correct += (  
                # FILL IN HERE  
                y_pred == y_trgt  
            ).sum().item()  
            n_total += batch_in.size(0)  
        val_accr = (n_correct/n_total)  
        model.train() # back to train mode  
    return val_accr  
print ("Done")
```

Calculating Predictions and Comparing Accuracy

# 5. Evaluation Function

## ✓ Evaluation Function

```
[ ] def func_eval(model,data_iter,device):  
    with torch.no_grad():  
        model.eval() # evaluate (affects Dropout and BN)  
        n_total,n_correct = 0,0  
        for batch_in,batch_out in data_iter:  
            y_trgt = batch_out.to(device)  
            model_pred = model(  
                # FILL IN HERE  
                batch_in.view(-1, 28*28).to(device)  
            )  
            _,y_pred = torch.max(model_pred.data,1)  
            n_correct += (  
                # FILL IN HERE  
                y_pred == y_trgt  
            ).sum().item()  
        # Returning to Training Mode  
        val_accr = (n_correct/n_total)  
        model.train() # back to train mode  
    return val_accr  
print ("Done")
```

# 5. Training model



## ▼ Train

```
[ ] print ("Start training.")
    M.init_param() # initialize parameters
    M.train()
    EPOCHS, print_every = 10, 1
    for epoch in range(EPOCHS):
        loss_val_sum = 0
        for batch_in, batch_out in train_iter:
            # Forward path
            y_pred = M.forward(batch_in.view(-1, 28*28).to(device))
            loss_out = loss(y_pred, batch_out.to(device))
            # Update
            # FILL IN HERE      # reset gradient
            optm.zero_grad()
            # FILL IN HERE      # backpropagate
            loss_out.backward()
            # FILL IN HERE      # optimizer update
            optm.step()
            loss_val_sum += loss_out
        loss_val_avg = loss_val_sum/len(train_iter)
        # Print
        if ((epoch%print_every)==0) or (epoch==(EPOCHS-1)):
            train_accr = func_eval(M, train_iter, device)
            test_accr = func_eval(M, test_iter, device)
            print ("epoch:[%d] loss:[%.3f] train_accr:[%.3f] test_accr:[%.3f]."%
                    (epoch, loss_val_avg, train_accr, test_accr))
    print ("Done")
```

# 5. Training model

▼ Train

```
[ ] print ("Start training.")  
    M.init_param() # initialize parameters  
    M.train()
```

Starting Training and Initializing Parameters

```
    EPOCHS, print_every = 10, 1  
    for epoch in range(EPOCHS):  
        for batch_in, batch_out in train_iter:  
            # Forward path  
            y_pred = M.forward(batch_in.view(-1, 28*28).to(device))  
            loss_out = loss(y_pred, batch_out.to(device))  
            # Update  
            # FILL IN HERE # reset gradient  
            optm.zero_grad()  
            # FILL IN HERE # backpropagate  
            loss_out.backward()  
            # FILL IN HERE # optimizer update  
            optm.step()  
            loss_val_sum += loss_out  
        loss_val_avg = loss_val_sum / len(train_iter)  
        # Print  
        if ((epoch % print_every) == 0) or (epoch == (EPOCHS - 1)):  
            train_accr = func_eval(M, train_iter, device)  
            test_accr = func_eval(M, test_iter, device)  
            print ("epoch: [%d] loss: [%.3f] train_accr: [%.3f] test_accr: [%.3f]."%  
                  (epoch, loss_val_avg, train_accr, test_accr))  
    print ("Done")
```

# 5. Training model

## Train

```
[ ] print ("Start training.")
    M.init_param() # initialize parameters
    M.train()
    EPOCHS, print_every = 10, 1
    for epoch in range(EPOCHS):
        loss_val_sum = 0
        for batch_in, batch_out in train_iter:
            # Forward path
            y_pred = M.forward(batch_in.view(-1, 28*28).to(device))
            loss_out = loss(y_pred, batch_out.to(device))
            # Update
            # FILL IN HERE # reset gradient
            # FILL IN HERE # backpropagate
            loss_out.backward()
            # FILL IN HERE # optimizer update
            optm.step()
            loss_val_sum += loss_out
        loss_val_avg = loss_val_sum / len(train_iter)
        # Print
        if ((epoch % print_every) == 0) or (epoch == (EPOCHS - 1)):
            train_accr = func_eval(M, train_iter, device)
            test_accr = func_eval(M, test_iter, device)
            print ("epoch: [%d] loss: [%.3f] train_accr: [%.3f] test_accr: [%.3f]."%
                    (epoch, loss_val_avg, train_accr, test_accr))
    print ("Done")
```

Epoch Loop and Batch Processing



# 5. Training model

## Train

```
[ ] print ("Start training.")
    M.init_param() # initialize parameters
    M.train()
    EPOCHS, print_every = 10, 1
    for epoch in range(EPOCHS):
        loss_val_sum = 0
        for batch_in, batch_out in train_iter:
            # Forward path
            y_pred = M.forward(batch_in.view(-1, 28*28).to(device))
            loss_out = loss(y_pred, batch_out.to(device))

            # Update
            # FILL IN HERE      # reset gradient
            optm.zero_grad()
            # FILL IN HERE      # backpropagate
            loss_out.backward()
            # FILL IN HERE      # optimizer update
            optm.step()
            loss_val_sum += loss_out
        loss_val_avg = loss_val_sum/len(train_iter)
        # Print

        Backpropagation and Weight Update-1):
        test_accr = func_eval(M, test_iter, device)
        print ("epoch:[%d] loss:[%.3f] train_accr:[%.3f] test_accr:[%.3f]."%
              (epoch, loss_val_avg, train_accr, test_accr))
    print ("Done")
```

# 6. Test model

## ▼ Test

```
[ ] n_sample = 25
    sample_indices = np.random.choice(len(mnist_test.targets), n_sample, replace=False)
    test_x = mnist_test.data[sample_indices]
    test_y = mnist_test.targets[sample_indices]
    with torch.no_grad():
        y_pred = M.forward(test_x.view(-1, 28*28).type(torch.float).to(device)/255.)
    y_pred = y_pred.argmax(axis=1)
    plt.figure(figsize=(10,10))
    for idx in range(n_sample):
        plt.subplot(5, 5, idx+1)
        plt.imshow(test_x[idx], cmap='gray')
        plt.axis('off')
        plt.title("Pred:%d, Label:%d"%(y_pred[idx],test_y[idx]))
    plt.show()
    print ("Done")
```

# 6. Test model

## ▼ Test

```
[ ] n_sample = 25
    sample_indices = np.random.choice(len(mnist_test.targets), n_sample, replace=False)
    test_x = mnist_test.data[sample_indices]
    test_y = mnist_test.targets[sample_indices]
    with torch.no_grad():
        y_pred = y_pred.argmax(axis=1)
        plt.figure(figsize=(10,10))
        for idx in range(n_sample):
            plt.subplot(5, 5, idx+1)
            plt.imshow(test_x[idx], cmap='gray')
            plt.axis('off')
            plt.title("Pred:%d, Label:%d"%(y_pred[idx],test_y[idx]))
        plt.show()
    print ("Done")
```

Selecting Test Samples

# 6. Test model

## ▼ Test

```
[ ] n_sample = 25
    sample_indices = np.random.choice(len(mnist_test.targets), n_sample, replace=False)
    test_x = mnist_test.data[sample_indices]
    test_y = mnist_test.targets[sample_indices]
    with torch.no_grad():
        y_pred = M.forward(test_x.view(-1, 28*28).type(torch.float).to(device)/255.)
        y_pred = y_pred.argmax(axis=1)
    plt.figure(figsize=(10,10))
    Model Prediction n_sample):
        plt.subplot(5, 5, idx+1)
        plt.imshow(test_x[idx], cmap='gray')
        plt.axis('off')
        plt.title("Pred:%d, Label:%d"%(y_pred[idx],test_y[idx]))
    plt.show()
    print ("Done")
```

# 6. Test model

## ▼ Test

```
[ ] n_sample = 25
    sample_indices = np.random.choice(len(mnist_test.targets), n_sample, replace=False)
    test_x = mnist_test.data[sample_indices]
    test_y = mnist_test.targets[sample_indices]
    with torch.no_grad():
        test_x.view(-1, 28*28).type(torch.float).to(device)/255.)
        y_pred = y_pred.argmax(axis=1)
        plt.figure(figsize=(10,10))
        for idx in range(n_sample):
            plt.subplot(5, 5, idx+1)
            plt.imshow(test_x[idx], cmap='gray')
            plt.axis('off')
            plt.title("Pred:%d, Label:%d"%(y_pred[idx],test_y[idx]))
        plt.show()
    print ("Done")
```

Visualizing the Predictions

# 6. Test model

## Test

```
[ ] n_sample = 25
    sample_indices = np.random.choice(mnist_test.data.shape[0], n_sample)
    test_x = mnist_test.data[sample_indices]
    test_y = mnist_test.targets[sample_indices]
    with torch.no_grad():
        y_pred = M.forward(test_x.view(-1, 1, 28, 28))
    y_pred = y_pred.argmax(axis=1)
    plt.figure(figsize=(10,10))
    for idx in range(n_sample):
        plt.subplot(5, 5, idx+1)
        plt.imshow(test_x[idx], cmap='gray')
        plt.axis('off')
        plt.title("Pred:%d, Label:%d" % (y_pred[idx], test_y[idx]))
    plt.show()
    print ("Done")
```

